

0000 0000 00 00000000000000 00
 00000000000000000000000000 000 0000000000

00000000 0000
 00000 000

00000000

0 0000000000 0000 00 000 0000 00000000

0 00000000 0 0000000000 0000000000 0 000000 0000 000000

000 0000

0 000 00000000 0 00000000000 00000 0000000000000

0 0000000000000000 0 0000 00000000000000

0 0000000000000 00 000000000 0 000000000000000 000000

000000000


0 00000000 0 00000000000 00000000 000 000000

0000000000000000 000000000000000 00000

0 000 00000000000 000 000000000000000000000

Blackwell Architecture Performance Comparison

Performance comparison of NVIDIA Blackwell architecture GPUs across different precision formats.



Blackwell architecture GPUs deliver exceptional performance across all precision formats, with FP8 showing the highest performance gains over previous generations.


RTX PRO 6000 is the only GPU in this comparison that supports FP8 performance.

For more detailed performance metrics and benchmarks, visit the NVIDIA website.

Blackwell architecture is designed for high-performance AI and data center workloads.

FP8 Performance

FP8 Performance Comparison





FP8 Performance Comparison

FP8 Performance Comparison

RTX PRO 6000 (a0001) — uniform

RTX PRO 6000 (a0001) — uniform




RTX PRO 6000 (a0001) — uniform

RTX PRO 6000 (a0001) — uniform

RTX PRO 6000 (a0001) — uniform

RTX PRO 6000 (a0001) — uniform

nvlink bandwidth




nvlink bandwidth is limited by the number of links between nodes. In a 4-way node-to-node configuration, the maximum bandwidth is 900 GB/s/dir. The achieved bandwidth for various operations is shown in the chart above.

The operations are ranked by their achieved bandwidth. The all_reduce operation achieves the highest bandwidth at 841 GB/s (93% of the NVLink max). The sendrecv operation achieves the lowest bandwidth at 666 GB/s (74% of the NVLink max).

The chart shows that the achieved bandwidth for all operations is significantly below the NVLink max, indicating that the operations are not fully utilizing the available bandwidth.

RDMA Performance

RDMA Performance Comparison: GDRDMA vs. sendrecv



RDMA Performance Comparison: GDRDMA vs. sendrecv

RDMA Performance Comparison: GDRDMA vs. sendrecv

RDMA Performance Comparison: GDRDMA vs. sendrecv


RDMA Performance Comparison: GDRDMA vs. sendrecv

RDMA Performance Comparison: GDRDMA vs. sendrecv

RDMA Performance Comparison: GDRDMA vs. sendrecv

RDMA Performance Comparison: GDRDMA vs. sendrecv

GPU bandwidth comparison



GPU bandwidth comparison results for various operations across three configurations: 1-node 2-GPU, 1-node 4-GPU, and 2-node 2-GPU. The chart shows that 1-node 4-GPU consistently achieves the highest bandwidth, followed by 2-node 2-GPU, and then 1-node 2-GPU. The scatter (unidir) operation shows the highest bandwidth for 1-node 4-GPU at 50.6 GB/s.

GPU bandwidth comparison results for various operations across three configurations: 1-node 2-GPU, 1-node 4-GPU, and 2-node 2-GPU. The chart shows that 1-node 4-GPU consistently achieves the highest bandwidth, followed by 2-node 2-GPU, and then 1-node 2-GPU. The scatter (unidir) operation shows the highest bandwidth for 1-node 4-GPU at 50.6 GB/s.

GPU bandwidth comparison results for various operations across three configurations: 1-node 2-GPU, 1-node 4-GPU, and 2-node 2-GPU. The chart shows that 1-node 4-GPU consistently achieves the highest bandwidth, followed by 2-node 2-GPU, and then 1-node 2-GPU. The scatter (unidir) operation shows the highest bandwidth for 1-node 4-GPU at 50.6 GB/s.

